

CN Assignment 3 Report

Srimant Mohanty

2021207

srimant21207@iiitd.ac.in

Plots

The experiment involved capturing packets exchanged between the client and the server on the server's interface and subsequently processing the acquired PCAP files using Python scripts. Two key metrics: the Average Throughput for each TCP flow (measured in bits per second) and the Average Latency for each TCP flow (measured in milliseconds) were computed to assess the performance of various TCP server programs that we developed as part of our experiments.

The following plots were obtained for 500, 1000, and 3000 simultaneous connections for each experiment.

Observations


Throughput

- The plots represent average throughput in bits per second and the average latency is milliseconds for each TCP flow during each scenario/experiment. The X-axis represents the different TCP flows which may be uniquely identified by a 4-tuple Source IP Address, Source Port Number, Destination IP address and Destination Port Number.
- The plots show that the average throughput values remain relatively consistent across all the 5 experiments but with an increasing number of TCP connections the throughput slightly decreases i.e. from 500 to 1000 to 3000 Connections.
This can be especially observed in cases of Select, Poll & Epoll System Calls.
- Also it maybe observed in some scenarios that throughput increases initially on increasing number of clients however after a point the network may become congested and throughput decreases



Latency

- The graphs plotted tend to show that average latency is consistent and constant with increasing number of concurrent connections and with different I/O multiplexing techniques.
- However, latency also tends to increase with an increasing number of clients in a few cases as the server has to manage more connections simultaneously, leading to longer response times for individual clients.
- I/O multiplexing techniques like select, epoll and poll are comparatively faster than traditional methods like forking and threading for handling concurrent clients.


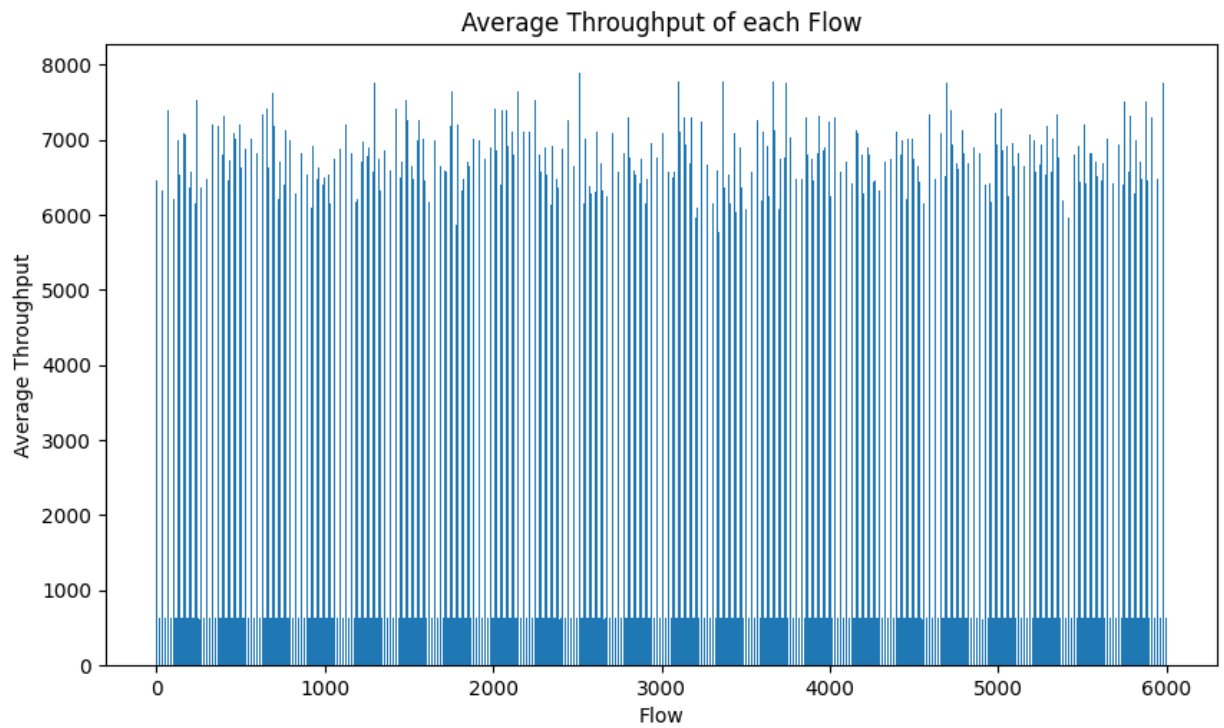
Fork 500 Connections





Fork 1000 Connections



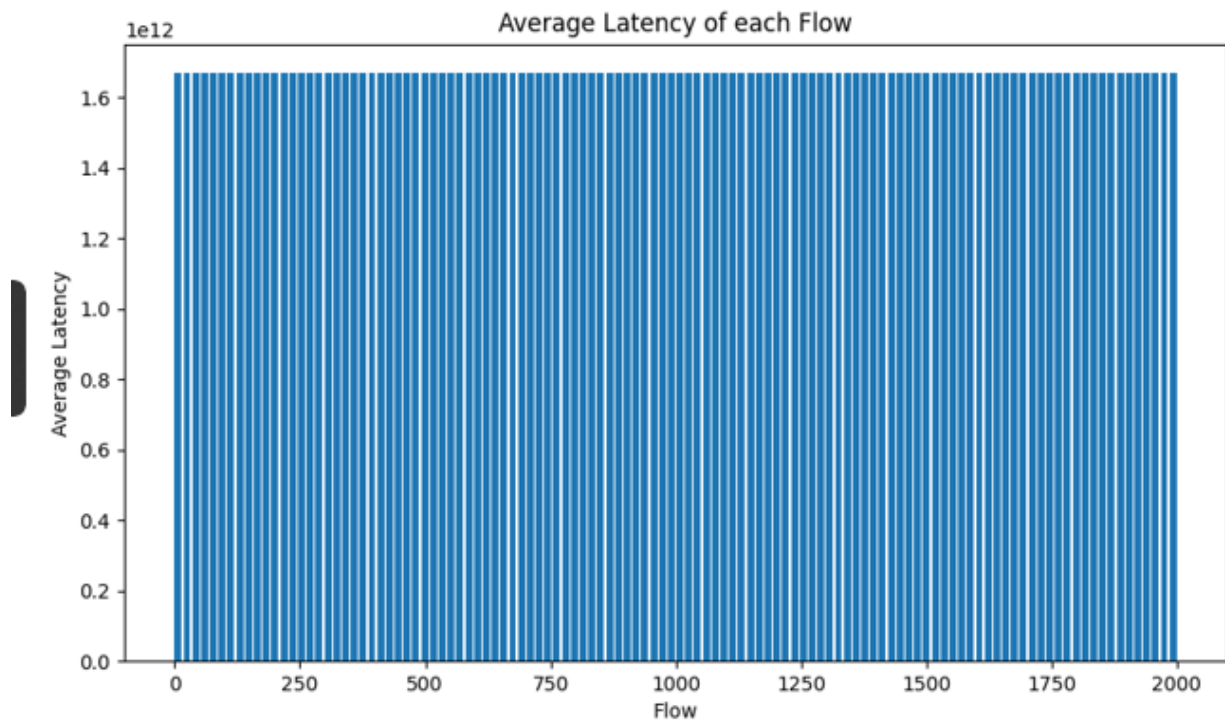
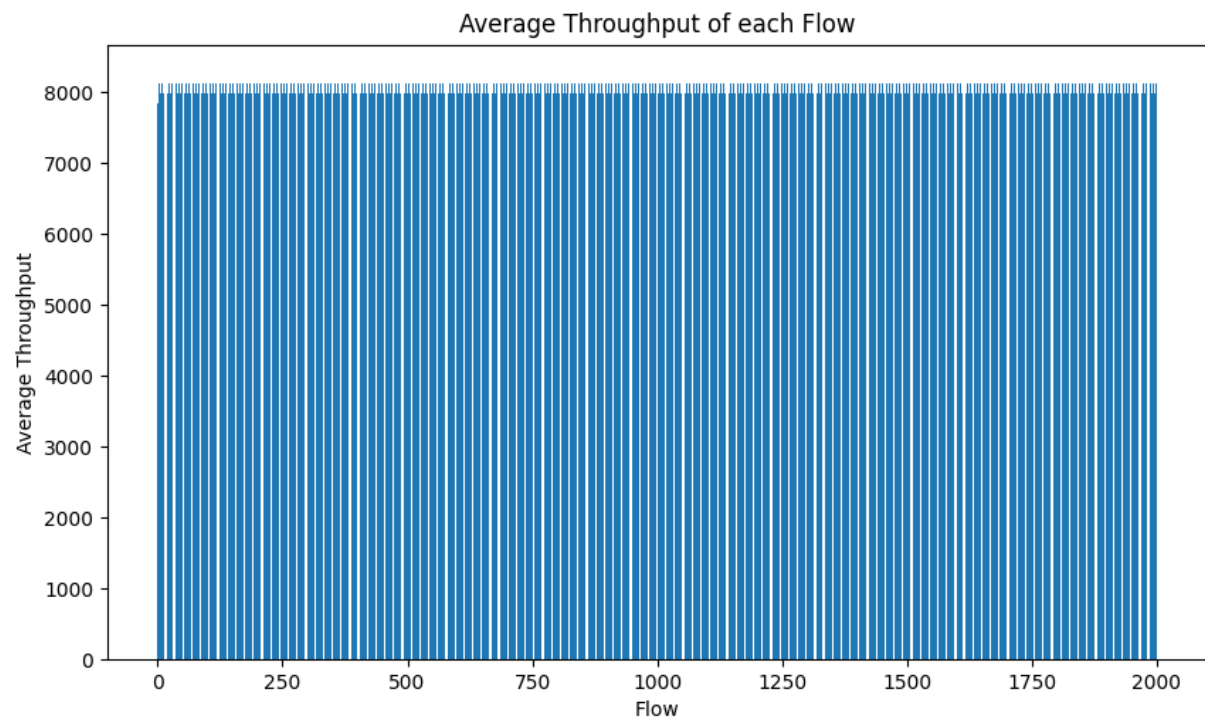
Fork 3000 Connections



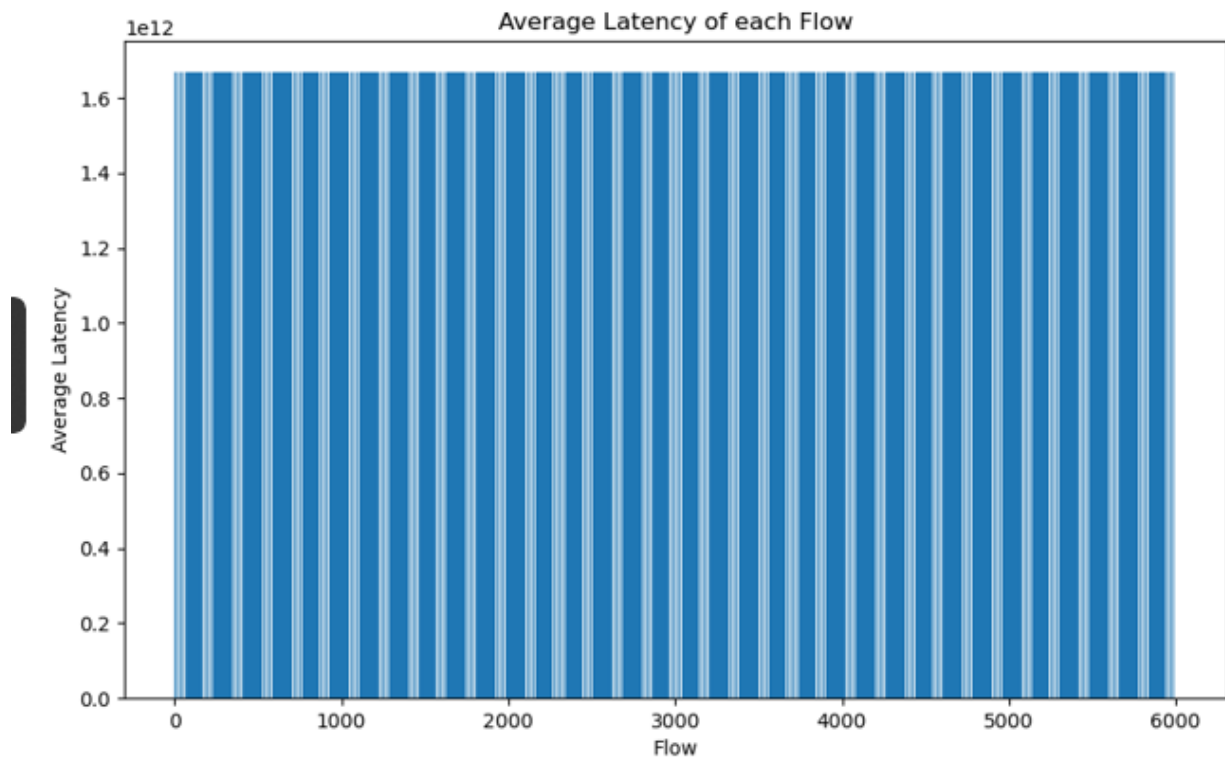
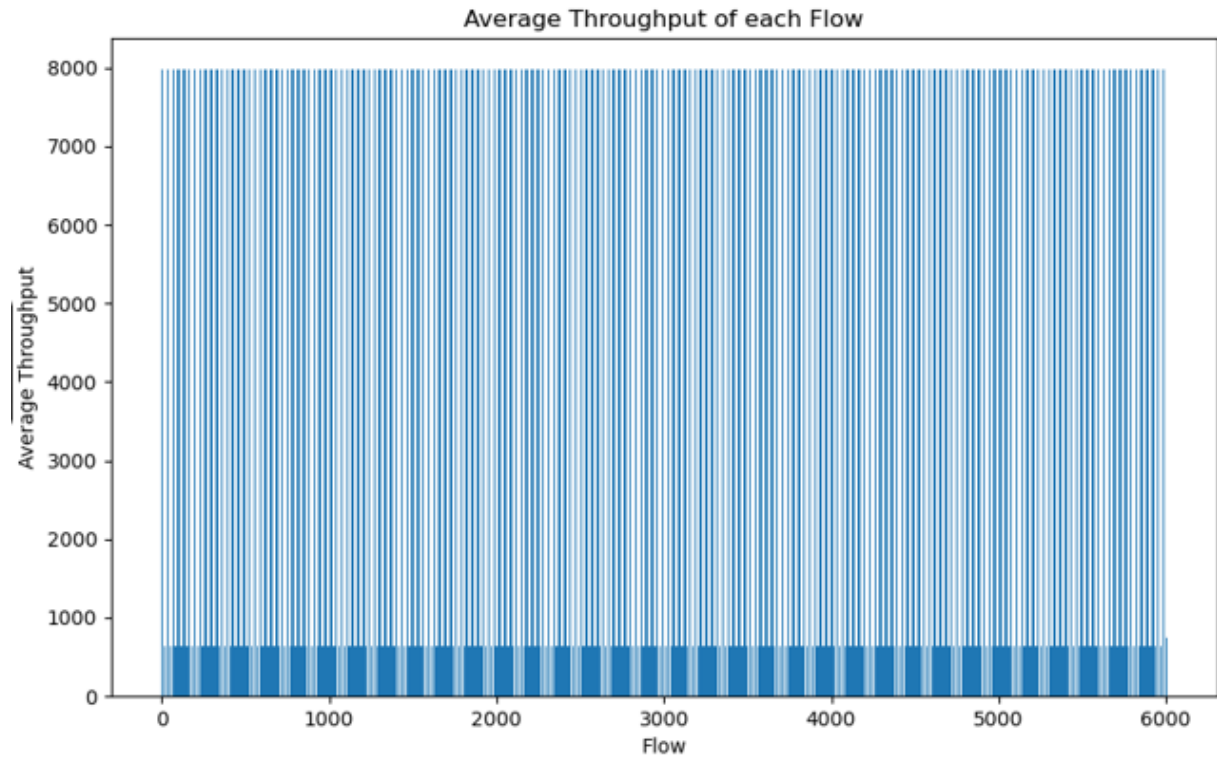
Thread 500 Connections



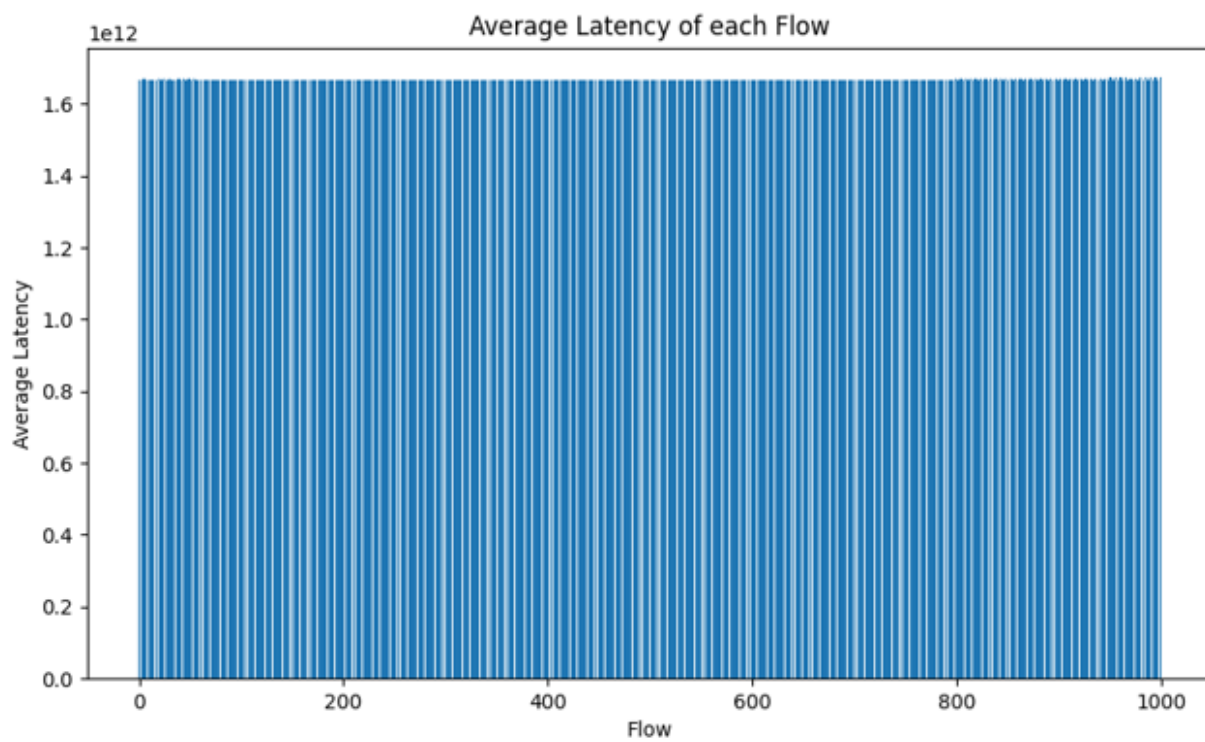
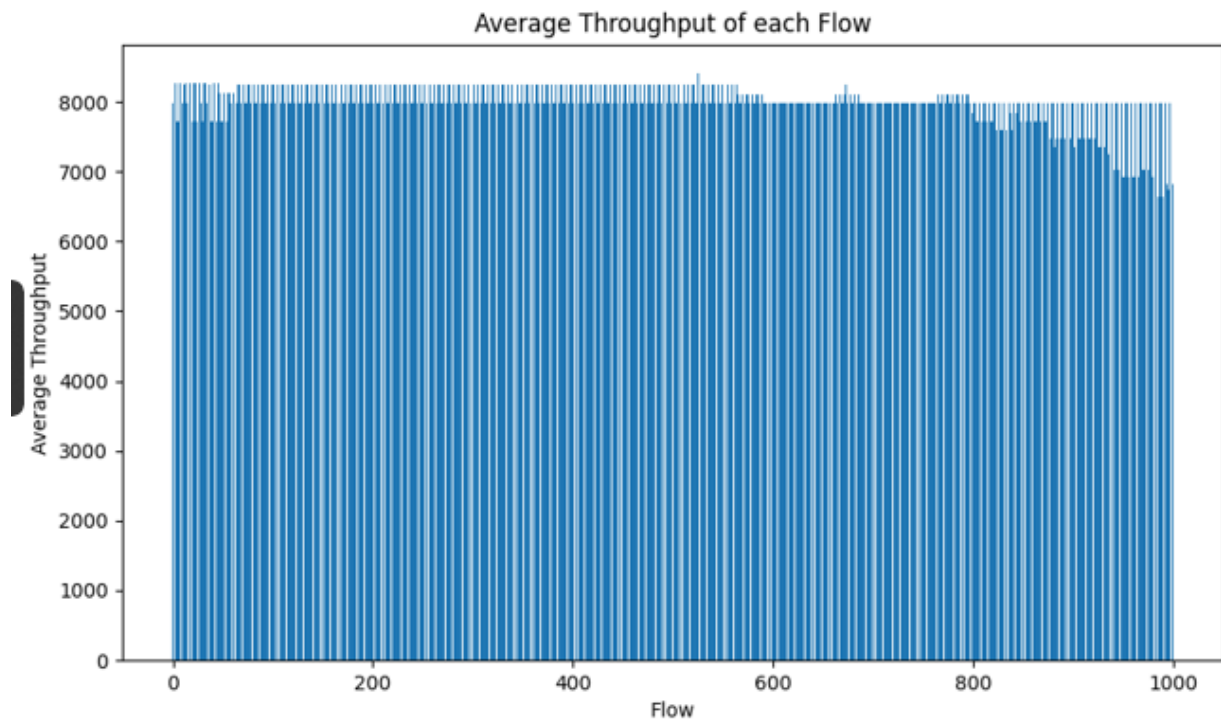
Thread 1000 Connections





Thread 3000 Connections





Select 500 Connections




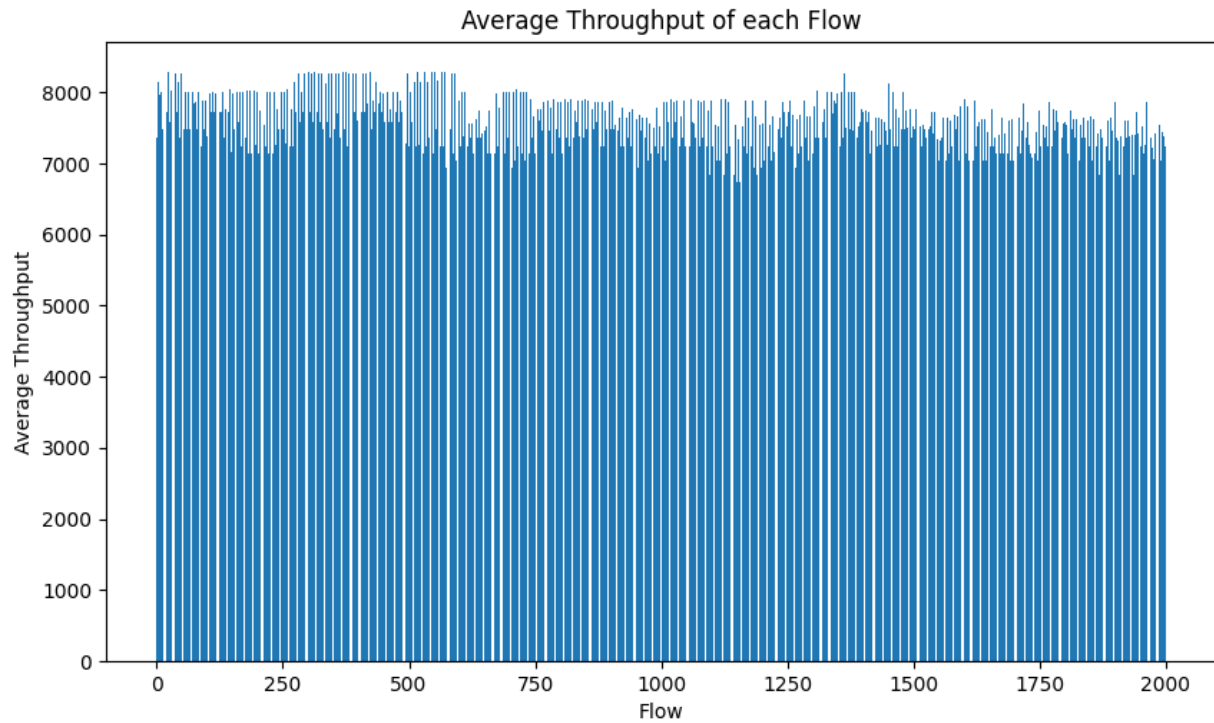
Select 1000 Connections



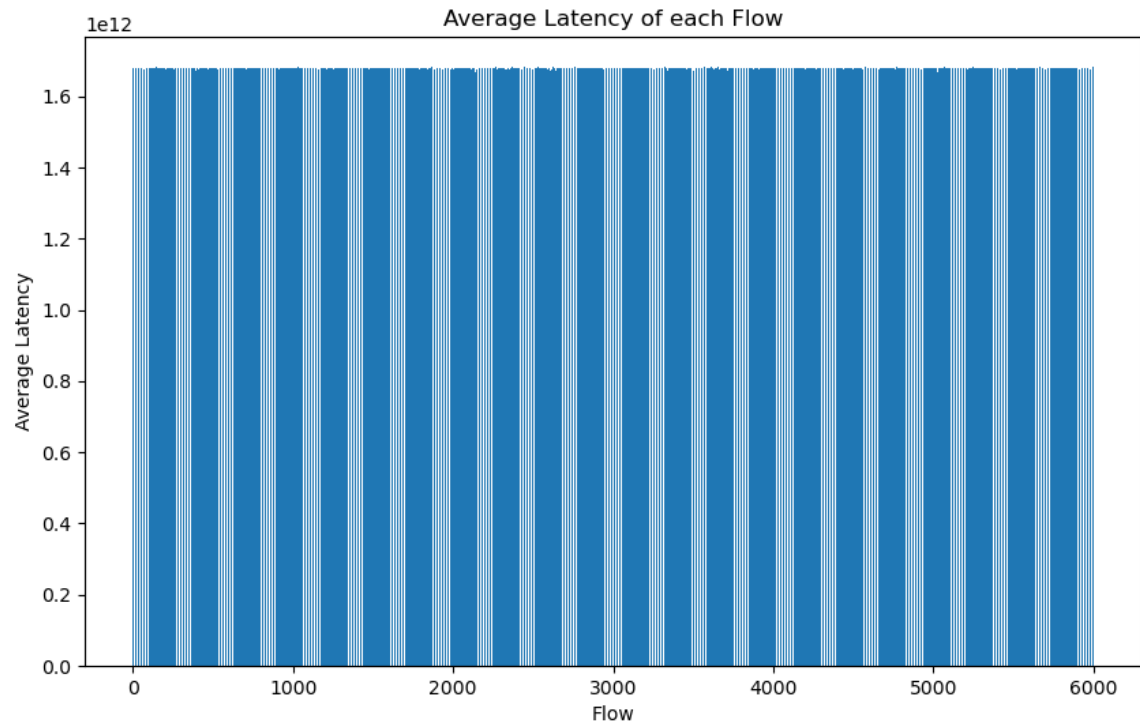

Poll 500 Connections




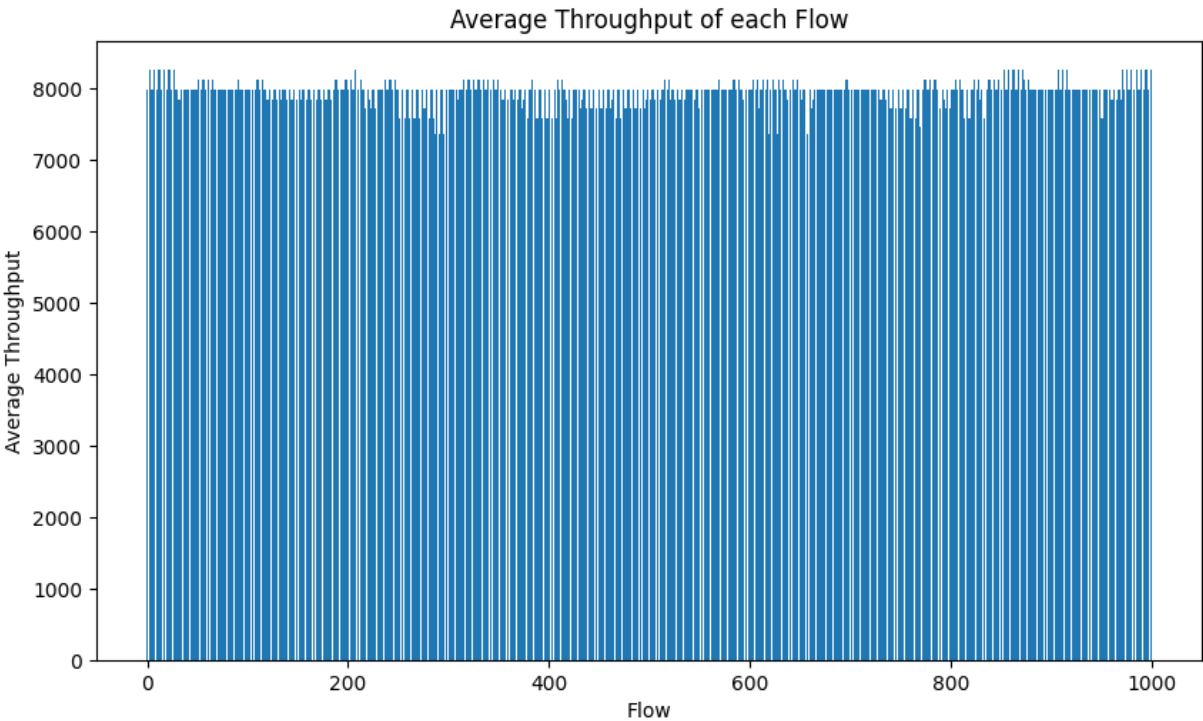
Poll 1000 Connections





Poll 3000 Connections



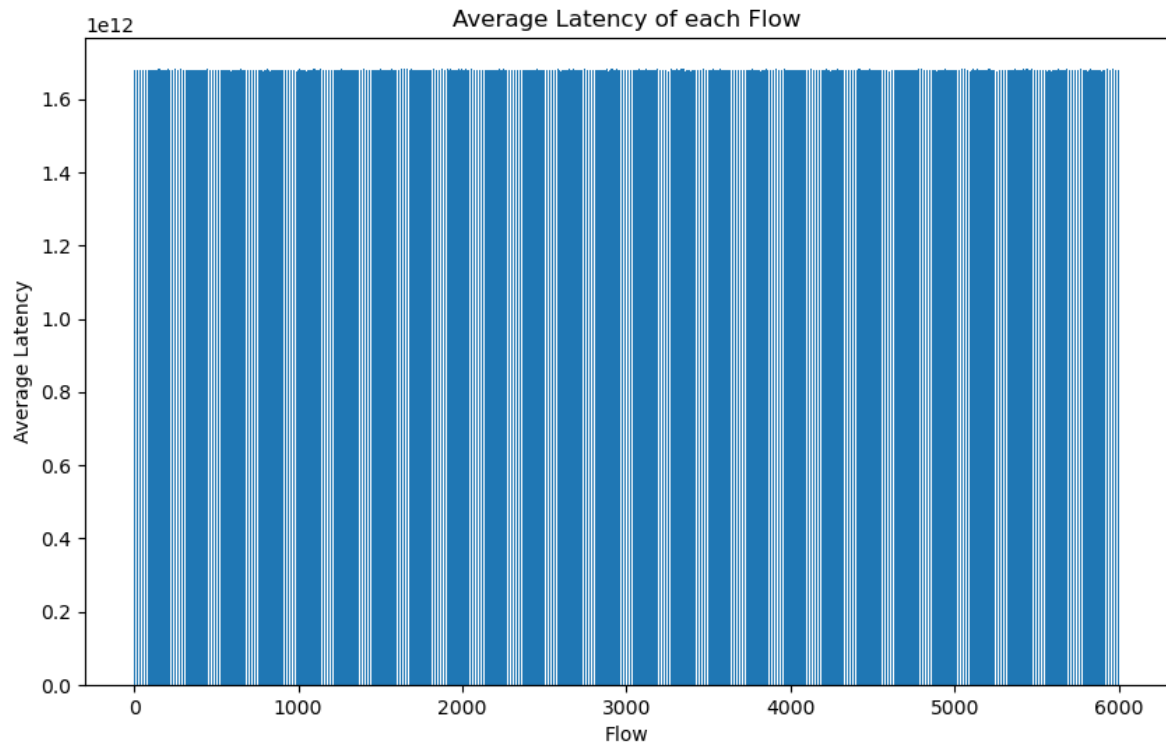

Epoll 500 Connections



Epoll 1000 Connections



Epoll 3000 Connections



Server process's CPU utilization

Process	CPU Utilization(%)
No Process	0.7%
Fork 500 Connections	11.7%
Fork 1000 Connections	49.8%
Fork 3000 Connections	66.5%
Thread 500 Connections	17.7%
Thread 1000 Connections	29.1%
Thread 3000 Connections	36.2%
Select 500 Connections	89%
Select 1000 Connections	103%
Poll 500 Connections	92.6%
Poll 1000 Connections	70.7%
Poll 3000 Connections	91%
EPoll 500 Connections	15.2%
EPoll 1000 Connections	42%
EPoll 3000 Connections	93.7%

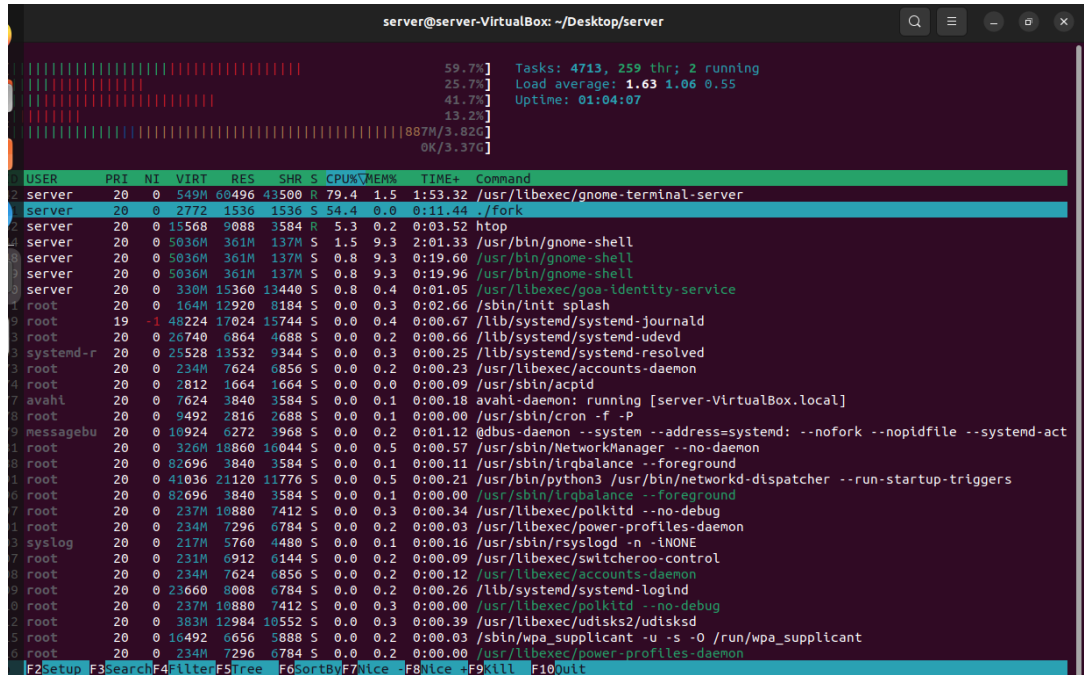
CPU utilization essentially represents the portion of a system's processing capacity that is currently in use.

Observation:

- A high CPU Utilization indicates that there is heavy load on the system due to the presence of multiple connections or multiple concurrent clients.
- We also observe that the CPU utilization values keep on varying during the period when a process is executing and the CPU Utilization % shown are those that occur more frequently or at random occurrences.
- The CPU Utilization % increases from 500 to 1000 and 3000 connections in every experiment indicating that more CPU power is required to serve multiple concurrent clients/requests.
- Different processes behave differently with respect to CPU Utilization time; for instance it drastically increases over time for a fork as more processes are being created in a fork. Similarly, I/O multiplexing techniques are more efficient as compared to fork and thread for CPU Utilization as their values are more or less stable.

- In some cases CPU Utilization may be more than 100 indicating that more than 1 core is being used.

CPU Utilization was captured in the following manner:-



Server process's memory utilization

For understanding the server's memory usage we observe the Maximum resident size which indicates the peak memory usage of the process.

We can see that the MRSS follows a stable and consistent pattern across all the experiments indicating that there is a stable and fixed memory utilization throughout all connections.

Process	Maximum resident size(K Bytes)
Fork 500 Connections	3328
Fork 1000 Connections	3456
Fork 3000 Connections	3456
Thread 500 Connections	3456
Thread 1000 Connections	3456
Thread 3000 Connections	3328
Select 500 Connections	3456
Select 1000 Connections	3456

Poll 500 Connections	3456
Poll 1000 Connections	3328
Poll 3000 Connections	3456
EPoll 500 Connections	3456
EPoll 1000 Connections	3456
EPoll 3000 Connections	3456

```

server@server-VirtualBox:~/Desktop/server$ /usr/bin/time -v 19605
/usr/bin/time: cannot run 19605: No such file or directory
Command exited with non-zero status 127
  Command being timed: "19605"
    User time (seconds): 0.00
    System time (seconds): 0.00
  Percent of CPU this job got: 87%
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:00.00
Average shared text size (kbytes): 0
Average unshared data size (kbytes): 0
Average stack size (kbytes): 0
Average total size (kbytes): 0
Maximum resident set size (kbytes): 1024
Average resident set size (kbytes): 0
Major (requiring I/O) page faults: 0
Minor (reclaiming a frame) page faults: 41
Voluntary context switches: 1
Involuntary context switches: 0
Swaps: 0
File system inputs: 0
File system outputs: 0
Socket messages sent: 0
Socket messages received: 0
Signals delivered: 0
Page size (bytes): 4096
Exit status: 127

```

References

<https://support.cloudways.com/en/articles/5120765-how-to-monitor-system-processes-using-htop-command>

<https://www.geeksforgeeks.org/using-htop-to-monitor-system-processes-on-linux/>

<https://www.softprayog.in/programming/io-multiplexing-select-poll-epoll-in-linux>